



(CSEN604) Data Bases II

Supervised by: Dr. Wael Abouelsaadat

Schema 1

Query 1:

Planning time and execution time with no index:

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Merge Full Join (cost=224.00..262.92 rows=66 width=63) (actual time=1.597..2.087 rows=66 loops=1)		
2	-> Seq Scan on student (cost=0.00..19.49 rows=66 width=23) (actual time=0.018..0.252 rows=66 loops=1)		
3	Filter: ((department)::text = 'CS1'::text)		
4	Rows Removed by Filter: 933		
5	-> Materialize (cost=224.00..242.61 rows=1 width=40) (actual time=1.575..1.806 rows=1 loops=1)		
6	-> Hash Join (cost=224.00..242.61 rows=1 width=40) (actual time=1.573..1.799 rows=1 loops=1)		
7	Hash Cond: (t.section_id = s.section_id)		
8	-> Seq Scan on takes t (cost=0.00..15.99 rows=999 width=12) (actual time=0.012..0.106 rows=999 loops=1)		
9	-> Hash (cost=223.98..223.98 rows=1 width=28) (actual time=1.549..1.550 rows=1 loops=1)		
10	Buckets: 1024 Batches: 1 Memory Usage: 9kB		
11	-> Seq Scan on section s (cost=0.00..223.98 rows=1 width=28) (actual time=0.007..1.543 rows=1 loops=1)		
12	Filter: ((semester = 1) AND (year = 2019))		
13	Rows Removed by Filter: 9998		
14	Planning Time: 0.206 ms		
15	Execution Time: 2.122 ms		

Planning and execution time using B+Trees on department:

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Merge Full Join (cost=228.66..255.92 rows=66 width=63) (actual time=1.060..1.271 rows=66 loops=1)		
2	-> Bitmap Heap Scan on student (cost=4.66..12.49 rows=66 width=23) (actual time=0.007..0.010 rows=66 loops=1)		
3	Recheck Cond: ((department)::text = 'CS1'::text)		
4	Heap Blocks: exact=7		
5	-> Bitmap Index Scan on index1 (cost=0.00..4.65 rows=66 width=0) (actual time=0.000..0.007 rows=66 loops=1)		
6	Index Cond: ((department)::text = 'CS1'::text)		
7	-> Materialize (cost=224.00..242.61 rows=1 width=40) (actual time=1.029..1.204 rows=1 loops=1)		
8	-> Hash Join (cost=224.00..242.61 rows=1 width=40) (actual time=1.020..1.193 rows=1 loops=1)		
9	Hash Cond: (t.section_id = s.section_id)		
10	-> Seq Scan on takes t (cost=0.00..15.99 rows=999 width=12) (actual time=0.000..0.106 rows=999 loops=1)		
11	-> Hash (cost=223.98..223.98 rows=1 width=28) (actual time=0.992..0.992 rows=1 loops=1)		
12	Buckets: 1024 Batches: 1 Memory Usage: 9kB		
13	-> Seq Scan on section s (cost=0.00..223.98 rows=1 width=28) (actual time=0.000..1.543 rows=1 loops=1)		
14	Filter: ((semester = 1) AND (year = 2019))		
15	Rows Removed by Filter: 9998		
16	Planning Time: 0.303 ms		
17	Execution Time: 1.327 ms		

Planning and execution time using B+Trees on year and semester:

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Merge Full Join (cost=8.32..47.24 rows=66 width=63) (actual time=0.058..0.288 rows=66 loops=1)		
2	-> Seq Scan on student (cost=0.00..19.49 rows=66 width=23) (actual time=0.017..0.114 rows=66 loops=1)		
3	Filter: ((department)::text = 'CS1'::text)		
4	Rows Removed by Filter: 933		
5	-> Materialize (cost=8.32..26.93 rows=1 width=40) (actual time=0.039..0.160 rows=1 loops=1)		
6	-> Hash Join (cost=8.32..26.93 rows=1 width=40) (actual time=0.035..0.154 rows=1 loops=1)		
7	Hash Cond: (t.section_id = s.section_id)		
8	-> Seq Scan on takes t (cost=0.00..15.99 rows=999 width=12) (actual time=0.010..0.071 rows=999 loops=1)		
9	-> Hash (cost=8.30..8.30 rows=1 width=28) (actual time=0.017..0.017 rows=1 loops=1)		
10	Buckets: 1024 Batches: 1 Memory Usage: 9kB		
11	-> Index Scan using query2 on section s (cost=0.29..8.30 rows=1 width=28) (actual time=0.000..0.000 rows=1 loops=1)		
12	Index Cond: ((semester = 1) AND (year = 2019))		
13	Planning Time: 0.229 ms		
14	Execution Time: 0.319 ms		

Planning and execution time using hash on semester:

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Merge Full Join (cost=8.03..46.96 rows=66 width=63) (actual time=0.052..0.292 rows=66 loops=1)		
2	-> Seq Scan on student (cost=0.00..19.49 rows=66 width=23) (actual time=0.017..0.126 rows=66 loops=1)		
3	Filter: ((department)::text = 'CS1'::text)		
4	Rows Removed by Filter: 933		
5	-> Materialize (cost=8.03..26.65 rows=1 width=40) (actual time=0.032..0.152 rows=1 loops=1)		
6	-> Hash Join (cost=8.03..26.64 rows=1 width=40) (actual time=0.031..0.149 rows=1 loops=1)		
7	Hash Cond: (t.section_id = s.section_id)		
8	-> Seq Scan on takes t (cost=0.00..15.99 rows=999 width=12) (actual time=0.010..0.067 rows=999 loops=1)		
9	-> Hash (cost=8.02..8.02 rows=1 width=28) (actual time=0.011..0.011 rows=1 loops=1)		
10	Buckets: 1024 Batches: 1 Memory Usage: 9kB		
11	-> Index Scan using query3 on section s (cost=0.00..8.02 rows=1 width=28) (actual time=0.000..0.000 rows=1 loops=1)		
12	Index Cond: (semester = 1)		
13	Filter: (year = 2019)		
14	Planning Time: 0.219 ms		
15	Execution Time: 0.324 ms		

Using hash on department:

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Merge Full Join (cost=228.51..255.77 rows=66 width=63) (actual time=0.737..0.887 rows=66 loops=1)		
2	-> Bitmap Heap Scan on student (cost=4.51..12.34 rows=66 width=23) (actual time=0.018..0.035 rows=66 loops=1)		
3	Recheck Cond: ((department)::text = 'CS1'::text)		
4	Heap Blocks: exact=7		
5	-> Bitmap Index Scan on query4 (cost=0.00..4.50 rows=66 width=0) (actual time=0.011..0.011 rows=66 loops=1)		
6	Index Cond: ((department)::text = 'CS1'::text)		
7	-> Materialize (cost=224.00..242.61 rows=1 width=40) (actual time=0.715..0.837 rows=1 loops=1)		
8	-> Hash Join (cost=224.00..242.61 rows=1 width=40) (actual time=0.711..0.830 rows=1 loops=1)		
9	Hash Cond: (t.section_id = s.section_id)		
10	-> Seq Scan on takes t (cost=0.00..15.99 rows=999 width=12) (actual time=0.011..0.064 rows=999 loops=1)		
11	-> Hash (cost=223.98..223.98 rows=1 width=28) (actual time=0.688..0.688 rows=1 loops=1)		
12	Buckets: 1024 Batches: 1 Memory Usage: 9kB		
13	-> Seq Scan on section s (cost=0.00..223.98 rows=1 width=28) (actual time=0.009..0.681 rows=1 loops=1)		
14	Filter: ((semester = 1) AND (year = 2019))		
15	Rows Removed by Filter: 9998		
16	Planning Time: 0.226 ms		
17	Execution Time: 0.932 ms		

Using hash on year:

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Merge Full Join (cost=224.00..262.92 rows=66 width=63) (actual time=1.065..1.404 rows=66 loops=1)		
2	-> Seq Scan on student (cost=0.00..19.49 rows=66 width=23) (actual time=0.025..0.163 rows=66 loops=1)		
3	Filter: ((department)::text = 'CS1')::text)		
4	Rows Removed by Filter: 933		
5	-> Materialize (cost=224.00..242.61 rows=1 width=40) (actual time=1.037..1.220 rows=1 loops=1)		
6	-> Hash Join (cost=224.00..242.61 rows=1 width=40) (actual time=1.033..1.213 rows=1 loops=1)		
7	Hash Cond: (t.section_id = s.section_id)		
8	-> Seq Scan on takes t (cost=0.00..15.99 rows=999 width=12) (actual time=0.017..0.100 rows=999 loops=1)		
9	-> Hash (cost=223.98..223.98 rows=1 width=28) (actual time=1.002..1.002 rows=1 loops=1)		
10	Buckets: 1024 Batches: 1 Memory Usage: 9kB		
11	-> Seq Scan on section s (cost=0.00..223.98 rows=1 width=28) (actual time=0.010..0.993 rows=1 loops=1)		
12	Filter: ((semester = 1) AND (year = 2019))		
13	Rows Removed by Filter: 9998		
14	Planning Time: 0.303 ms		
15	Execution Time: 1.448 ms		

Planning and execution time using mixed indices B+tree on year and semester and hash on department:

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Merge Full Join (cost=12.83..40.09 rows=66 width=63) (actual time=0.060..0.208 rows=66 loops=1)		
2	-> Bitmap Heap Scan on student (cost=4.51..12.34 rows=66 width=23) (actual time=0.019..0.036 rows=66 loops=1)		
3	Recheck Cond: ((department)::text = 'CS1')::text)		
4	Heap Blocks: exact=7		
5	-> Bitmap Index Scan on query4 (cost=0.00..4.50 rows=66 width=0) (actual time=0.013..0.013 rows=66 loops=1)		
6	Index Cond: ((department)::text = 'CS1')::text)		
7	-> Materialize (cost=8.32..26.93 rows=1 width=40) (actual time=0.039..0.159 rows=1 loops=1)		
8	-> Hash Join (cost=8.32..26.93 rows=1 width=40) (actual time=0.035..0.154 rows=1 loops=1)		
9	Hash Cond: (t.section_id = s.section_id)		
10	-> Seq Scan on takes t (cost=0.00..15.99 rows=999 width=12) (actual time=0.012..0.061 rows=999 loops=1)		
11	-> Hash (cost=8.30..8.30 rows=1 width=28) (actual time=0.016..0.016 rows=1 loops=1)		
12	Buckets: 1024 Batches: 1 Memory Usage: 9kB		
13	-> Index Scan using query2 on section s (cost=0.29..8.30 rows=1 width=28) (actual time=0.000..0.000 rows=1 loops=1)		
14	Index Cond: ((semester = 1) AND (year = 2019))		
15	Planning Time: 0.231 ms		
16	Execution Time: 0.251 ms		

Report Summary:

By checking the results shown above we can deduce that the execution time significantly improved when using indices. Best option would be mix between B+tree on year and semester and hash on department so that all columns we are performing a query on are indexed. B+ tree runs faster when on year and semester or even on of them only as B+tree is better used on numerical data. Hash is insignificantly slower than B+tree.

Without index: 2.122 ms

With B+tree on department: 1.327 ms

With B+tree on year and semester: 0.319 ms

With hash on semester: 0.324ms

With hash on department: 0.932ms

With hash on year: 1.448ms

With mixed indices B+tree on year and semester and hash on department: 0.251ms

Schema 2

Query 2

Scenario 1(No index)

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	HashAggregate (cost=423.68..425.56 rows=188 width=4) (actual time=10.985..11.055 rows=599 loops=1)		
2	Group Key: project.pnumber		
3	Batches: 1 Memory Usage: 89kB		
4	-> Seq Scan on project (cost=409.46..423.21 rows=188 width=4) (actual time=10.616..10.798 rows=599 loops=1)		
5	Filter: ((hashed SubPlan 1) OR (hashed SubPlan 2))		
6	SubPlan 1		
7	-> Nested Loop (cost=0.15..158.20 rows=1 width=4) (actual time=0.036..7.823 rows=17371 loops=1)		
8	-> Nested Loop (cost=0.15..143.20 rows=1 width=0) (actual time=0.031..4.829 rows=29 loops=1)		
9	-> Seq Scan on employee e (cost=0.00..105.58 rows=9 width=8) (actual time=0.010..1.162 rows=4999 loops=1)		
10	Filter: (lname = 'employee1':bpchar)		
11	Rows Removed by Filter: 2		
12	-> Index Scan using department_pkey on department d (cost=0.15..4.17 rows=1 width=8) (actual time=0.000..0.00...		
13	Index Cond: (dnumber = e.dno)		
14	Filter: (e.ssn = mgr.ssn)		
15	-> Seq Scan on project project_1 (cost=0.00..12.50 rows=250 width=4) (actual time=0.002..0.046 rows=599 loops=29)		
16	SubPlan 2		
17	-> Nested Loop (cost=4.06..251.19 rows=27 width=4) (never executed)		
18	-> Seq Scan on employee (cost=0.00..105.58 rows=9 width=4) (never executed)		
19	Filter: (lname = 'employee1':bpchar)		
20	-> Bitmap Heap Scan on works_on (cost=4.06..15.89 rows=29 width=8) (never executed)		
21	Recheck Cond: (essn = employee.ssn)		
22	-> Bitmap Index Scan on works_on_pkey (cost=0.00..4.06 rows=29 width=0) (never executed)		
23	Index Cond: (essn = employee.ssn)		
24	Planning Time: 0.888 ms		
25	Execution Time: 11.218 ms		

Scenario 2(B+ Tree index on Ssn in Employee)

“Used B+ index on ssn because it is accessed in every loop that happens in the query, thus shortening the time”

1	HashAggregate (cost=543.70..545.58 rows=188 width=4) (actual time=10.729..10.802 rows=599 loops=1)
2	Group Key: project.pnumber
3	Batches: 1 Memory Usage: 89kB
4	-> Seq Scan on project (cost=529.48..543.23 rows=188 width=4) (actual time=10.369..10.527 rows=59...
5	Filter: ((hashed SubPlan 1) OR (hashed SubPlan 2))
6	SubPlan 1
7	-> Nested Loop (cost=26.25..185.89 rows=1 width=4) (actual time=0.042..6.249 rows=17371 loops...
8	-> Hash Join (cost=26.25..170.89 rows=1 width=0) (actual time=0.035..1.871 rows=29 loops=1)
9	Hash Cond: ((e.dno = d.dnumber) AND (e.ssn = d.mgr_ssn))
10	-> Seq Scan on employee e (cost=0.00..144.51 rows=25 width=8) (actual time=0.007..1.041 ...
11	Filter: (lname = 'employee1':bpchar)
12	Rows Removed by Filter: 2
13	-> Hash (cost=16.50..16.50 rows=650 width=8) (actual time=0.017..0.018 rows=29 loops=1)
14	Buckets: 1024 Batches: 1 Memory Usage: 10kB
15	-> Seq Scan on department d (cost=0.00..16.50 rows=650 width=8) (actual time=0.005.....
16	-> Seq Scan on project project_1 (cost=0.00..12.50 rows=250 width=4) (actual time=0.002..0.06...

17	SubPlan 2
18	-> Nested Loop (cost=0.28..343.51 rows=29 width=4) (never executed)
19	-> Seq Scan on employee (cost=0.00..144.51 rows=25 width=4) (never executed)
20	Filter: (lname = 'employee1':bpchar)
21	-> Index Only Scan using works_on_pkey on works_on (cost=0.28..7.67 rows=29 width=8) (never executed)
22	Index Cond: (essn = employee.ssn)
23	Heap Fetches: 0
24	Planning Time: 0.881 ms
25	Execution Time: 10.980 ms

Scenario 3(Hash index on Ssn in Employee)

“Used Hash index on ssn because it is accessed in every loop that happens in the query, thus shortening the time”

1	HashAggregate (cost=543.70..545.58 rows=188 width=4) (actual time=10.679..10.748 rows=599 loops=1)
2	Group Key: project.pnumber
3	Batches: 1 Memory Usage: 89kB
4	-> Seq Scan on project (cost=529.48..543.23 rows=188 width=4) (actual time=10.318..10.490 rows=599 loops=1)
5	Filter: ((hashed SubPlan 1) OR (hashed SubPlan 2))
6	SubPlan 1
7	-> Nested Loop (cost=26.25..185.89 rows=1 width=4) (actual time=0.093..6.289 rows=17371 loops=1)
8	-> Hash Join (cost=26.25..170.89 rows=1 width=0) (actual time=0.074..1.852 rows=29 loops=1)
9	Hash Cond: ((e.dno = d.dnumber) AND (e.ssn = d.mgr_ssn))
10	-> Seq Scan on employee e (cost=0.00..144.51 rows=25 width=8) (actual time=0.019..1.023 rows=4999 loops=1)
11	Filter: (lname = 'employee1':bpchar)
12	Rows Removed by Filter: 2
13	-> Hash (cost=16.50..16.50 rows=650 width=8) (actual time=0.034..0.035 rows=29 loops=1)
14	Buckets: 1024 Batches: 1 Memory Usage: 10kB
15	-> Seq Scan on department d (cost=0.00..16.50 rows=650 width=8) (actual time=0.009..0.018 rows=29 loops=1)
16	-> Seq Scan on project project_1 (cost=0.00..12.50 rows=250 width=4) (actual time=0.003..0.071 rows=599 loops=1)

17	SubPlan 2
18	-> Nested Loop (cost=0.28..343.51 rows=29 width=4) (never executed)
19	-> Seq Scan on employee (cost=0.00..144.51 rows=25 width=4) (never executed)
20	Filter: (lname = 'employee1':bpchar)
21	-> Index Only Scan using works_on_pkey on works_on (cost=0.28..7.67 rows=29 width=8) (never executed)
22	Index Cond: (essn = employee.ssn)
23	Heap Fetches: 0
24	Planning Time: 0.723 ms
25	Execution Time: 10.964 ms

Scenario 4(Mixed index [Hash on Ssn, B+ on Dnumber on Department])

“I used here the a hash index on ssn because it is accessed at everyloop and also the dnumber is accessed many times so I used B+ index with it thus the time decreased drastically compared to the 3 previous runs”

1	HashAggregate (cost=519.18..521.06 rows=188 width=4) (actual time=8.189..8.356 rows=599 loops=1)
2	Group Key: project.pnumber
3	Batches: 1 Memory Usage: 89kB
4	-> Seq Scan on project (cost=504.96..518.71 rows=188 width=4) (actual time=7.140..7.675 rows=599 loops=1)
5	Filter: ((hashed SubPlan 1) OR (hashed SubPlan 2))
6	SubPlan 1
7	-> Nested Loop (cost=1.73..161.37 rows=1 width=4) (actual time=0.047..4.933 rows=17371 loops=1)
8	-> Hash Join (cost=1.73..146.37 rows=1 width=0) (actual time=0.039..2.531 rows=29 loops=1)
9	Hash Cond: ((e.dno = d.dnumber) AND (e.ssn = d.mgr_ssn))
10	-> Seq Scan on employee e (cost=0.00..144.51 rows=25 width=8) (actual time=0.012..1.567 rows=4999 loops=1)
11	Filter: (lname = 'employee1':bpchar)
12	Rows Removed by Filter: 2
13	-> Hash (cost=1.29..1.29 rows=29 width=8) (actual time=0.019..0.021 rows=29 loops=1)
14	Buckets: 1024 Batches: 1 Memory Usage: 10kB
15	-> Seq Scan on department d (cost=0.00..1.29 rows=29 width=8) (actual time=0.006..0.011 rows=29 loops=1)
16	-> Seq Scan on project project_1 (cost=0.00..12.50 rows=250 width=4) (actual time=0.001..0.038 rows=599 loops=1)

17	SubPlan 2
18	-> Nested Loop (cost=0.28..343.51 rows=29 width=4) (never executed)
19	-> Seq Scan on employee (cost=0.00..144.51 rows=25 width=4) (never executed)
20	Filter: (lname = 'employee1':bpchar)
21	-> Index Only Scan using works_on_pkey on works_on (cost=0.28..7.67 rows=29 width=8) (never executed)
22	Index Cond: (essn = employee.ssn)
23	Heap Fetches: 0
24	Planning Time: 3.273 ms
25	Execution Time: 8.738 ms

Query 3

Scenario 1(No index)

1	Seq Scan on employee (cost=0.00..361966.86 rows=2500 width=168) (actual time=1.405..4.287 rows=167...
2	Filter: (SubPlan 1)
3	Rows Removed by Filter: 3329
4	SubPlan 1
5	-> Materialize (cost=0.00..144.64 rows=25 width=4) (actual time=0.000..0.000 rows=1 loops=5001)
6	-> Seq Scan on employee employee_1 (cost=0.00..144.51 rows=25 width=4) (actual time=0.008..1.3...
7	Filter: (dno = 5)
8	Rows Removed by Filter: 5000
9	Planning Time: 0.161 ms
10	Execution Time: 4.393 ms

Scenario 2(B+ Tree index on Dnumber in Department)

“All the retrieved events should have a salary of those who have the dnumber=5, thus ordering by using the B+ index, the dnumber will surely make the execution faster”

1	Seq Scan on employee (cost=0.00..361966.86 rows=2500 width=168) (actual time=1.126..3.480 rows=1672 loops=1)
2	Filter: (SubPlan 1)
3	Rows Removed by Filter: 3329
4	SubPlan 1
5	-> Materialize (cost=0.00..144.64 rows=25 width=4) (actual time=0.000..0.000 rows=1 loops=5001)
6	-> Seq Scan on employee employee_1 (cost=0.00..144.51 rows=25 width=4) (actual time=0.009..1.102 rows=1 loops=1)
7	Filter: (dno = 5)
8	Rows Removed by Filter: 5000
9	Planning Time: 0.151 ms
10	Execution Time: 3.575 ms

Scenario 3(Hash index on Salary in Employee)

“All the retrieved employees should have salary>than all the salaries of those who have d=no5, thus ordering by using the Hash index, the salary will surely make the execution faster”

1	Seq Scan on employee (cost=0.00..361966.86 rows=2500 width=168) (actual time=0.884..2.888 rows=1672 loops=1)
2	Filter: (SubPlan 1)
3	Rows Removed by Filter: 3329
4	SubPlan 1
5	-> Materialize (cost=0.00..144.64 rows=25 width=4) (actual time=0.000..0.000 rows=1 loops=5001)
6	-> Seq Scan on employee employee_1 (cost=0.00..144.51 rows=25 width=4) (actual time=0.007..0.866 rows=1 loops=1)
7	Filter: (dno = 5)
8	Rows Removed by Filter: 5000
9	Planning Time: 0.403 ms
10	Execution Time: 2.960 ms

Scenario 4(Mixed index[B+ on dnumber,Hash on salary])

“Using both indexes on dnumber the B+ and on the salary using the Hash resulted in a shorter execution time than the original query”

1	Seq Scan on employee (cost=0.00..361966.86 rows=2500 width=168) (actual time=1.060..3.019 rows=1672 loops=1)
2	Filter: (SubPlan 1)
3	Rows Removed by Filter: 3329
4	SubPlan 1
5	-> Materialize (cost=0.00..144.64 rows=25 width=4) (actual time=0.000..0.000 rows=1 loops=5001)
6	-> Seq Scan on employee employee_1 (cost=0.00..144.51 rows=25 width=4) (actual time=0.013..1.029 rows=1 loops=1)
7	Filter: (dno = 5)
8	Rows Removed by Filter: 5000
9	Planning Time: 0.531 ms
10	Execution Time: 3.115 ms

Query 4

Scenario 1(No index)

1	Seq Scan on employee e (cost=0.00..216512.78 rows=2500 width=168) (actual time=0.067..6763.901 rows=4999 loops=1)
2	Filter: (SubPlan 1)
3	Rows Removed by Filter: 2
4	SubPlan 1
5	-> Seq Scan on dependent d (cost=0.00..81.64 rows=1956 width=4) (actual time=0.470..1.117 rows=1251 loops=5001)
6	Filter: ((e.fname <> dependent_name) AND (e.sex <> sex))
7	Rows Removed by Filter: 1250
8	Planning Time: 12.303 ms
9	Execution Time: 6765.118 ms

Scenario 2(B+ Tree index on fname in Employee)

“fname is always compared with the dependent name so using the B+ index on the fname will make the execution time less”

1	Seq Scan on employee e (cost=0.00..216512.78 rows=2500 width=168) (actual time=0.018..1486.254 rows=4999 loops=1)
2	Filter: (SubPlan 1)
3	Rows Removed by Filter: 2
4	SubPlan 1
5	-> Seq Scan on dependent d (cost=0.00..81.64 rows=1956 width=4) (actual time=0.105..0.248 rows=1251 loops=5001)
6	Filter: ((e.fname <> dependent_name) AND (e.sex <> sex))
7	Rows Removed by Filter: 1250
8	Planning Time: 1.009 ms
9	Execution Time: 1486.514 ms

Scenario 3(Hash index on sex in Employee)

“the sex of the employee is always compared so using also a hash index with will make the execution tim less”

1	Seq Scan on employee e (cost=0.00..216512.78 rows=2500 width=168) (actual time=0.030..1438.155 rows=4999 loops=1)
2	Filter: (SubPlan 1)
3	Rows Removed by Filter: 2
4	SubPlan 1
5	-> Seq Scan on dependent d (cost=0.00..81.64 rows=1956 width=4) (actual time=0.101..0.239 rows=1251 loops=5001)
6	Filter: ((e.fname <> dependent_name) AND (e.sex <> sex))
7	Rows Removed by Filter: 1250
8	Planning Time: 0.379 ms
9	Execution Time: 1438.380 ms

Scenario 4(Mixed index[B+ on fname, Hash on sex])

“Using 2 index increased the planning time but decreased the overall execution time and resulted in the best performance”

1	Seq Scan on employee e (cost=0.00..216512.78 rows=2500 width=168) (actual time=0.019..1427.008 rows=4999 loops=1)
2	Filter: (SubPlan 1)
3	Rows Removed by Filter: 2
4	SubPlan 1
5	-> Seq Scan on dependent d (cost=0.00..81.64 rows=1956 width=4) (actual time=0.101..0.237 rows=1251 loops=5001)
6	Filter: ((e.fname <> dependent_name) AND (e.sex <> sex))
7	Rows Removed by Filter: 1250
8	Planning Time: 0.626 ms
9	Execution Time: 1427.198 ms

Query 5

Scenario 1(No index)

1	Hash Join (cost=118.99..291.94 rows=2500 width=168) (actual time=1.908..3.337 rows=4999 loops=1)
2	Hash Cond: (employee.ssn = dependent.essn)
3	-> Seq Scan on employee (cost=0.00..132.01 rows=5001 width=172) (actual time=0.008..0.504 rows=5001 loops=1)
4	-> Hash (cost=116.49..116.49 rows=200 width=4) (actual time=1.895..1.896 rows=4999 loops=1)
5	Buckets: 8192 (originally 1024) Batches: 1 (originally 1) Memory Usage: 240kB
6	-> HashAggregate (cost=114.49..116.49 rows=200 width=4) (actual time=1.125..1.467 rows=4999 loops=1)
7	Group Key: dependent.essn
8	Batches: 1 Memory Usage: 481kB
9	-> Seq Scan on dependent (cost=0.00..101.99 rows=4999 width=4) (actual time=0.004..0.330 rows=4999 loops=1)
10	Planning Time: 0.430 ms
11	Execution Time: 3.758 ms

Scenario 2(B+ Tree index on Ssn in Employee)

“The query only check on the ssn in employee and compare it with essn in dependent so creating a B+ index on ssn normally decreased the time”

1	Hash Join (cost=81.20..248.32 rows=1976 width=168) (actual time=2.075..3.087 rows=4999 loops=1)
2	Hash Cond: (employee.ssn = dependent.essn)
3	-> Seq Scan on employee (cost=0.00..132.01 rows=5001 width=172) (actual time=0.006..0.387 rows=5001 loops=1)
4	-> Hash (cost=78.70..78.70 rows=200 width=4) (actual time=2.060..2.061 rows=4999 loops=1)
5	Buckets: 8192 (originally 1024) Batches: 1 (originally 1) Memory Usage: 240kB
6	-> HashAggregate (cost=76.70..78.70 rows=200 width=4) (actual time=1.277..1.665 rows=4999 loops=1)
7	Group Key: dependent.essn
8	Batches: 1 Memory Usage: 481kB
9	-> Seq Scan on dependent (cost=0.00..71.76 rows=1976 width=4) (actual time=0.004..0.377 rows=4999 loops=1)
10	Planning Time: 0.328 ms
11	Execution Time: 3.288 ms

Scenario 3(Hash index on essn in dependent)

“The query only check on the ssn in employee and compare it with essn in dependent so creating a B+ index on essn normally decreased the time”

1	Hash Join (cost=81.20..248.32 rows=1976 width=168) (actual time=2.330..3.264 rows=4999 loops=1)
2	Hash Cond: (employee.ssn = dependent.essn)
3	-> Seq Scan on employee (cost=0.00..132.01 rows=5001 width=172) (actual time=0.012..0.317 rows=5001 loops=1)
4	-> Hash (cost=78.70..78.70 rows=200 width=4) (actual time=2.310..2.311 rows=4999 loops=1)
5	Buckets: 8192 (originally 1024) Batches: 1 (originally 1) Memory Usage: 240kB
6	-> HashAggregate (cost=76.70..78.70 rows=200 width=4) (actual time=1.278..1.754 rows=4999 loops=1)
7	Group Key: dependent.essn
8	Batches: 1 Memory Usage: 481kB
9	-> Seq Scan on dependent (cost=0.00..71.76 rows=1976 width=4) (actual time=0.007..0.377 rows=4999 loops=1)
10	Planning Time: 0.204 ms
11	Execution Time: 3.451 ms

Scenario 4(Mixed index[B+ on SSn,Hash on essn])

“The query only check on the ssn in employee and compare it with essn in dependent so creating a B+ index on ssn and essn resulted in the least execution time”

1	Hash Join (cost=118.99..291.94 rows=2500 width=168) (actual time=1.845..3.020 rows=4999 loops=1)
2	Hash Cond: (employee.ssn = dependent.essn)
3	-> Seq Scan on employee (cost=0.00..132.01 rows=5001 width=172) (actual time=0.013..0.416 rows=5001 loops=1)
4	-> Hash (cost=116.49..116.49 rows=200 width=4) (actual time=1.826..1.828 rows=4999 loops=1)
5	Buckets: 8192 (originally 1024) Batches: 1 (originally 1) Memory Usage: 240kB
6	-> HashAggregate (cost=114.49..116.49 rows=200 width=4) (actual time=1.151..1.456 rows=4999 loops=1)
7	Group Key: dependent.essn
8	Batches: 1 Memory Usage: 481kB
9	-> Seq Scan on dependent (cost=0.00..101.99 rows=4999 width=4) (actual time=0.005..0.326 rows=4999 loops=1)
10	Planning Time: 0.663 ms
11	Execution Time: 3.260 ms

Query 6

Scenario 1(No index)

1	GroupAggregate (cost=159.51..318.02 rows=1 width=12) (actual time=2.696..2.697 rows=1 loops=1)
2	Group Key: department.dnumber
3	InitPlan 1 (returns \$0)
4	-> HashAggregate (cost=157.01..159.51 rows=67 width=4) (actual time=1.911..2.052 rows=1 loops=1)
5	Group Key: employee_1.dno
6	Filter: (count(*) > 2)
7	Batches: 1 Memory Usage: 737kB
8	Rows Removed by Filter: 4998
9	-> Seq Scan on employee employee_1 (cost=0.00..132.01 rows=5001 width=4) (actual time=0.009..0.486 rows=5001 loops=1)
10	-> Nested Loop (cost=0.00..158.46 rows=8 width=4) (actual time=2.078..2.692 rows=1 loops=1)
11	-> Seq Scan on department (cost=0.00..1.36 rows=1 width=4) (actual time=2.063..2.064 rows=1 loops=1)
12	Filter: (dnumber = \$0)
13	Rows Removed by Filter: 28
14	-> Seq Scan on employee (cost=0.00..157.01 rows=8 width=4) (actual time=0.012..0.624 rows=1 loops=1)
15	Filter: ((salary > 40000) AND (dno = \$0))
16	Rows Removed by Filter: 5000
17	Planning Time: 0.403 ms
18	Execution Time: 2.820 ms

Scenario 2(B+ Tree index on Dnumber on department)

“In the nested query the dno is always compared so having a B+ index on it decreased the time taken to execute it”

1	GroupAggregate (cost=159.51..318.02 rows=1 width=12) (actual time=2.431..2.432 rows=1 loops=1)	Read-only column
2	Group Key: department.dnumber	
3	InitPlan 1 (returns \$0)	
4	-> HashAggregate (cost=157.01..159.51 rows=67 width=4) (actual time=1.710..1.842 rows=1 loops=1)	
5	Group Key: employee_1.dno	
6	Filter: (count(*) > 2)	
7	Batches: 1 Memory Usage: 737kB	
8	Rows Removed by Filter: 4998	
9	-> Seq Scan on employee employee_1 (cost=0.00..132.01 rows=5001 width=4) (actual time=0.010..0.362 rows=5001 loops=1)	
10	-> Nested Loop (cost=0.00..158.46 rows=8 width=4) (actual time=1.862..2.429 rows=1 loops=1)	
11	-> Seq Scan on department (cost=0.00..1.36 rows=1 width=4) (actual time=1.853..1.856 rows=1 loops=1)	
12	Filter: (dnumber = \$0)	
13	Rows Removed by Filter: 28	
14	-> Seq Scan on employee (cost=0.00..157.01 rows=8 width=4) (actual time=0.007..0.571 rows=1 loops=1)	
15	Filter: ((salary > 40000) AND (dno = \$0))	
16	Rows Removed by Filter: 5000	
17	Planning Time: 0.116 ms	
18	Execution Time: 2.575 ms	

Scenario 3(Hash index on Salary in Employee)

“There is a condition where the salary should always be greater than 40000 thus having a hash index on it will decrease the execution time”

1	GroupAggregate (cost=159.51..318.02 rows=1 width=12) (actual time=2.204..2.255 rows=1 loops=1)
2	Group Key: department.dnumber
3	InitPlan 1 (returns \$0)
4	-> HashAggregate (cost=157.01..159.51 rows=67 width=4) (actual time=1.638..1.786 rows=1 loops=1)
5	Group Key: employee_1.dno
6	Filter: (count(*) > 2)
7	Batches: 1 Memory Usage: 737kB
8	Rows Removed by Filter: 4998
9	-> Seq Scan on employee employee_1 (cost=0.00..132.01 rows=5001 width=4) (actual time=0.009..0.400 row...
10	-> Nested Loop (cost=0.00..158.46 rows=8 width=4) (actual time=1.817..2.252 rows=1 loops=1)
11	-> Seq Scan on department (cost=0.00..1.36 rows=1 width=4) (actual time=1.794..1.796 rows=1 loops=1)
12	Filter: (dnumber = \$0)
13	Rows Removed by Filter: 28
14	-> Seq Scan on employee (cost=0.00..157.01 rows=8 width=4) (actual time=0.022..0.455 rows=1 loops=1)
15	Filter: ((salary > 40000) AND (dno = \$0))
16	Rows Removed by Filter: 5000
17	Planning Time: 0.953 ms
18	Execution Time: 2.388 ms

Scenario 4(Mixed index [Hash on Salary,B+ on dnumber])

“Having both indexes resulted in the best performance in the execution time”

1	GroupAggregate (cost=159.51..318.02 rows=1 width=12) (actual time=2.083..2.083 rows=1 loops=1)
2	Group Key: department.dnumber
3	InitPlan 1 (returns \$0)
4	-> HashAggregate (cost=157.01..159.51 rows=67 width=4) (actual time=1.507..1.645 rows=1 loops=1)
5	Group Key: employee_1.dno
6	Filter: (count(*) > 2)
7	Batches: 1 Memory Usage: 737kB
8	Rows Removed by Filter: 4998
9	-> Seq Scan on employee employee_1 (cost=0.00..132.01 rows=5001 width=4) (actual time=0.004..0.279 rows=5001 loops=1)
10	-> Nested Loop (cost=0.00..158.46 rows=8 width=4) (actual time=1.665..2.081 rows=1 loops=1)
11	-> Seq Scan on department (cost=0.00..1.36 rows=1 width=4) (actual time=1.657..1.658 rows=1 loops=1)
12	Filter: (dnumber = \$0)
13	Rows Removed by Filter: 28
14	-> Seq Scan on employee (cost=0.00..157.01 rows=8 width=4) (actual time=0.007..0.421 rows=1 loops=1)
15	Filter: ((salary > 40000) AND (dno = \$0))
16	Rows Removed by Filter: 5000
17	Planning Time: 0.547 ms
18	Execution Time: 2.206 ms

Schema 3

Query 7:

First of all, I analyzed both tables (sailors, reserves) to collect statistics to help us getting accurate estimates building the plan.

	schemaname	tablename	attname	inherited	null_frac	avg_width	n_distinct	most_common_vals	most_common_freqs	histogram_bounds
	name	name	name	boolean	real	integer	real	anyarray	real[]	anyarray
1	public	boat	bid	false	0	4	-1	[null]	[null]	(1,30,60,90,120,150,180,...)
2	public	boat	bname	false	0	21	-1	[null]	[null]	(*Boat1, ..., *Boat10,...)
3	public	boat	color	false	0	11	1	(*Red, ...)	[1]	[null]


*“select * from pg_stats where tablename = ‘boat’ ”*

	schemaname	tablename	attname	inherited	null_frac	avg_width	n_distinct	most_common_vals	most_common_freqs	histogram_bounds	correlation	most_common_elems
	name	name	name	boolean	real	integer	real	anyarray	real[]	anyarray	real	anyarray
1	public	reserves	day	false	0	4	1	(1906-07-23)	(1)	[null]	1	[null]
2	public	reserves	bid	false	0	4	-0.19886667	(355,1733,2999,1101,1202,...)	(.00066666666,0.0006,0.0006)	(1,32,61,90,121,152,178,...)	-0.011462337	[null]
3	public	reserves	sid	false	0	4	-0.485	(1469,1268,1327,2814,377,...)	(.0033333333,0.0033333333)	(1,102,192,292,383,474,5,...)	0.005622192	[null]

*“select * from pg_stats where tablename = ‘reserves’ ”*

Then I modified how PostgreSQL consider plans by disabling nested loops

“SET enable_nestloop=false”



pgAdmin 4

File Object Tools Help

Dashboard Properties SQL Statistics Dependencies Dependents schema3/postgres@PostgreSQL 13 *

schema3/postgres@PostgreSQL 13

Query Editor Query History

```

1 select s.sname
2 from sailors s
3 where
4 s.sid in( select r.sid
5 from reserves r
6 where r.bid = 103 );

```

Data Output Explain Messages Notifications

Graphical Analysis Statistics

#	Node	Rows			
		Rows X	Actual	Plan	Loops
1.	→ Hash Inner Join (cost=269.56..458.2 rows=5 width=21) (rows=6 loops=1) Hash Cond: (s.sid = r.sid)	11.2	6	5	1
2.	→ Seq Scan on sailors as s (cost=0..165 rows=9000 width=25) (rows=9000 loops=1)	11	9000	9000	1
3.	→ Hash (cost=269.5..269.5 rows=5 width=4) (rows=6 loops=1) Buckets: 1024 Batches: 1 Memory Usage: 9 kB	11.2	6	5	1
4.	→ Seq Scan on reserves as r (cost=0..269.5 rows=5 width=4) (rows=6 loops=1) Filter: (bid = 103) Rows Removed by Filter: 14994	11.2	6	5	1

"Planning Time: 0.352 ms"

"Execution Time: 6.027 ms"

Scenario 1 (No index):

1	explain analyze
2	select s.sname
3	from sailors s
4	where
5	s.sid in(select r.sid
6	from reserves r
7	where r.bid = 103);
8	

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Nested Loop (cost=0.29..311.01 rows=5 width=21) (actual time=0.025..0.855 ro...		
2	-> Seq Scan on reserves r (cost=0.00..269.50 rows=5 width=4) (actual time=0...		
3	Filter: (bid = 103)		
4	Rows Removed by Filter: 14994		
5	-> Index Scan using sailors_pkey on sailors s (cost=0.29..8.30 rows=1 width=2...		
6	Index Cond: (sid = r.sid)		
7	Planning Time: 1.067 ms		
8	Execution Time: 0.871 ms		

Scenario 2 (B+ tree index):

1	explain analyze
2	select s.sname
3	from sailors s
4	where
5	s.sid in(select r.sid
6	from reserves r
7	where r.bid = 103);
8	
9	

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Nested Loop (cost=0.29..291.01 rows=5 width=21) (actual time=0.032..0.870 rows=6 loops=1)		
2	-> Seq Scan on reserves r (cost=0.00..269.50 rows=5 width=4) (actual time=0.018..0.807 rows=6 loops=1)		
3	Filter: (bid = 103)		
4	Rows Removed by Filter: 14994		
5	-> Index Only Scan using idx on sailors s (cost=0.29..4.30 rows=1 width=25) (actual time=0.010..0.010 rows=1 loops=6)		
6	Index Cond: (sid = r.sid)		
7	Heap Fetches: 0		
8	Planning Time: 2.769 ms		
9	Execution Time: 0.886 ms		

Scenario 3 (Hash index):

1	explain analyze select s.sname
2	from sailors s
3	where
4	s.sid in(select r.sid
5	from reserves r
6	where r.bid = 103);

Data Output	Explain	Messages	Notifications
<div> <div>QUERY PLAN</div> <div> <div>text</div> <div> 1 Nested Loop (cost=4.04..60.49 rows=5 width=21) (actual time=0.305..0.326 rows=6 loops=1) 2 -> Bitmap Heap Scan on reserves r (cost=4.04..20.40 rows=5 width=4) (actual time=0.297..0.303 rows=6 loops=1) 3 Recheck Cond: (bid = 103) 4 Heap Blocks: exact=6 5 -> Bitmap Index Scan on idx4 (cost=0.00..4.04 rows=5 width=0) (actual time=0.291..0.291 rows=6 loops=1) 6 Index Cond: (bid = 103) 7 -> Index Scan using idx2 on sailors s (cost=0.00..8.02 rows=1 width=25) (actual time=0.003..0.003 rows=1 loops=6) 8 Index Cond: (sid = r.sid) 9 Planning Time: 3.469 ms 10 Execution Time: 0.362 ms </div> </div> </div>			

Execution time improved drastically. As I created hash index on r.bid since hash indexes work really well with exact value queries in addition to an extra index on s.sid.

Scenario 4 (mixed indexes):

1	explain analyze
2	select s.sname
3	from sailors s
4	where
5	s.sid in(select r.sid
6	from reserves r
7	where r.bid = 103);

Data Output	Explain	Messages	Notifications
<div> <div>QUERY PLAN</div> <div> <div>text</div> <div> 1 Nested Loop (cost=4.32..41.91 rows=5 width=21) (actual time=0.026..0.079 rows=6 loops=1) 2 -> Bitmap Heap Scan on reserves r (cost=4.04..20.40 rows=5 width=4) (actual time=0.011..0.017 rows=6 loops=1) 3 Recheck Cond: (bid = 103) 4 Heap Blocks: exact=6 5 -> Bitmap Index Scan on idx4 (cost=0.00..4.04 rows=5 width=0) (actual time=0.008..0.008 rows=6 loops=1) 6 Index Cond: (bid = 103) 7 -> Index Only Scan using idx0 on sailors s (cost=0.29..4.30 rows=1 width=25) (actual time=0.010..0.010 rows=1 loops=6) 8 Index Cond: (sid = r.sid) 9 Heap Fetches: 0 10 Planning Time: 2.805 ms 11 Execution Time: 0.105 ms </div> </div> </div>			

Execution time improved drastically. As each index help in a specific way in optimizing the query.

Query 8:

Scenario 1 (no index):

1	<code>explain analyze</code>
2	<code>select s.sname</code>
3	<code>from sailors s</code>
Data Output Explain Messages Notifications	
QUERY PLAN	
text	
1	Nested Loop (cost=358.24..359.77 rows=5 width=21) (actual time=4.282..11.843 rows=5123 loops=1)
2	→ HashAggregate (cost=357.96..358.01 rows=5 width=4) (actual time=4.276..5.077 rows=5123 loops=1)
3	Group Key: r.sid
4	Batches: 1 Memory Usage: 489kB
5	→ Hash Join (cost=64.70..357.95 rows=5 width=4) (actual time=0.523..2.941 rows=7569 loops=1)
6	Hash Cond: (r.bid = b.bid)
7	→ Seq Scan on reserves r (cost=0.00..249.28 rows=16728 width=8) (actual time=0.012..0.792 rows=15000 loops=1)
8	→ Hash (cost=64.69..64.69 rows=1 width=4) (actual time=0.505..0.506 rows=1500 loops=1)
9	Buckets: 2048 (originally 1024) Batches: 1 (originally 1) Memory Usage: 69kB
10	→ Seq Scan on boat b (cost=0.00..64.69 rows=1 width=4) (actual time=0.008..0.336 rows=1500 loops=1)
11	Filter: (color = 'red':bpchar)
12	Rows Removed by Filter: 1500
13	→ Index Scan using sailors_pkey on sailors s (cost=0.29..0.35 rows=1 width=25) (actual time=0.001..0.001 rows=1 loops=5123)
14	Index Cond: (sid = r.sid)
15	Planning Time: 0.256 ms
16	Execution Time: 12.061 ms

Scenario 2 (B+ tree index):

Query Editor Query History	
1	<code>explain analyze select s.sname</code>
2	<code>from sailors s</code>
3	<code>where s.sid in (select r.sid</code>
Data Output Explain Messages Notifications	
QUERY PLAN	
text	
1	Hash Semi Join (cost=658.28..1111.43 rows=7296 width=21) (actual time=5.856..8.239 rows=5123 loops=1)
2	Hash Cond: (s.sid = r.sid)
3	→ Index Scan using sailors_pkey on sailors s (cost=0.29..321.29 rows=9000 width=25) (actual time=0.011..1.061 rows=9000 loops=1)
4	→ Hash (cost=564.25..564.25 rows=7500 width=4) (actual time=5.829..5.831 rows=7569 loops=1)
5	Buckets: 8192 Batches: 1 Memory Usage: 331kB
6	→ Hash Join (cost=127.81..564.25 rows=7500 width=4) (actual time=0.526..4.984 rows=7569 loops=1)
7	Hash Cond: (r.bid = b.bid)
8	→ Index Only Scan using idx2 on reserves r (cost=0.29..397.29 rows=15000 width=8) (actual time=0.007..2.779 rows=15000 loops=1)
9	Heap Fetches: 0
10	→ Hash (cost=108.78..108.78 rows=1500 width=4) (actual time=0.514..0.515 rows=1500 loops=1)
11	Buckets: 2048 Batches: 1 Memory Usage: 69kB
12	→ Index Only Scan using idx3 on boat b (cost=0.28..108.78 rows=1500 width=4) (actual time=0.010..0.381 rows=1500 loops=1)
13	Filter: (color = 'red':bpchar)
14	Rows Removed by Filter: 1500
15	Heap Fetches: 0
16	Planning Time: 0.282 ms
17	Execution Time: 8.469 ms

To force Postgres to use the created indexes I used “set enable_seqscan=false”, “set enable_indexscan=true”

Scenario 3 (Hash index):

1	explain analyze
2	select s.sname
3	from sailors s

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Hash Semi Join (cost=704.88..1158.03 rows=7296 width=21) (actual time=5.534..7.5...		
2	Hash Cond: (s.sid = r.sid)		
3	-> Index Scan using sailors_pkey on sailors s (cost=0.29..321.29 rows=9000 width=2...		
4	-> Hash (cost=610.84..610.84 rows=7500 width=4) (actual time=5.469..5.471 rows=...		
5	Buckets: 8192 Batches: 1 Memory Usage: 331kB		
6	-> Hash Join (cost=122.41..610.84 rows=7500 width=4) (actual time=0.541..4.39...		
7	Hash Cond: (r.bid = b.bid)		
8	-> Index Only Scan using reserves_pkey on reserves r (cost=0.29..449.29 rows...		
9	Heap Fetches: 0		
10	-> Hash (cost=103.38..103.38 rows=1500 width=4) (actual time=0.504..0.505 ...		
11	Buckets: 2048 Batches: 1 Memory Usage: 69kB		
12	-> Bitmap Heap Scan on boat b (cost=59.63..103.38 rows=1500 width=4) (...)		
13	Recheck Cond: (color = 'red':bpchar)		
14	Heap Blocks: exact=13		
15	-> Bitmap Index Scan on idx11 (cost=0.00..59.25 rows=1500 width=0) (...)		
16	Index Cond: (color = 'red':bpchar)		
17	Planning Time: 2.749 ms		
18	Execution Time: 7.850 ms		

Scenario 4 (mixed indexes):

1	explain analyze
2	select s.sname
3	from sailors s

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Hash Semi Join (cost=652.88..1106.03 rows=7296 width=21) (actual time=3.907..6.1...		
2	Hash Cond: (s.sid = r.sid)		
3	-> Index Scan using sailors_pkey on sailors s (cost=0.29..321.29 rows=9000 width=2...		
4	-> Hash (cost=558.84..558.84 rows=7500 width=4) (actual time=3.874..3.876 rows=...		
5	Buckets: 8192 Batches: 1 Memory Usage: 331kB		
6	-> Hash Join (cost=122.41..558.84 rows=7500 width=4) (actual time=0.382..3.04...		
7	Hash Cond: (r.bid = b.bid)		
8	-> Index Only Scan using idx2 on reserves r (cost=0.29..397.29 rows=15000 wi...		
9	Heap Fetches: 0		
10	-> Hash (cost=103.38..103.38 rows=1500 width=4) (actual time=0.362..0.363 ...		
11	Buckets: 2048 Batches: 1 Memory Usage: 69kB		
12	-> Bitmap Heap Scan on boat b (cost=59.63..103.38 rows=1500 width=4) (...)		
13	Recheck Cond: (color = 'red':bpchar)		
14	Heap Blocks: exact=13		
15	-> Bitmap Index Scan on idx11 (cost=0.00..59.25 rows=1500 width=0) (...)		
16	Index Cond: (color = 'red':bpchar)		
17	Planning Time: 0.286 ms		
18	Execution Time: 6.315 ms		

Query 9:

Scenario 1 (no index):

1	explain analyze select s.sname
2	from sailors s, reserves r, boat b
3	where
Data Output Explain Messages Notifications	
QUERY PLAN	
text	
1	Hash Semi Join (cost=1102.38..1482.73 rows=6250 width=21) (actual time=9.706..18.805 rows=7569 loops=1)
2	Hash Cond: (s.sid = s2.sid)
3	→ Hash Join (cost=358.75..649.88 rows=7500 width=29) (actual time=2.979..9.931 rows=7569 loops=1)
4	Hash Cond: (r.sid = s.sid)
5	→ Hash Join (cost=81.25..352.68 rows=7500 width=4) (actual time=0.619..4.540 rows=7569 loops=1)
6	Hash Cond: (r.bid = b.bid)
7	→ Seq Scan on reserves r (cost=0.00..232.00 rows=15000 width=8) (actual time=0.014..0.980 rows=15000 loops=1)
8	→ Hash (cost=62.50..62.50 rows=1500 width=4) (actual time=0.596..0.597 rows=1500 loops=1)
9	Buckets: 2048 Batches: 1 Memory Usage: 69kB
10	→ Seq Scan on boat b (cost=0.00..62.50 rows=1500 width=4) (actual time=0.007..0.333 rows=1500 loops=1)
11	Filter: (color = 'red'::bpchar)
12	Rows Removed by Filter: 1500
13	→ Hash (cost=165.00..165.00 rows=9000 width=25) (actual time=2.311..2.311 rows=9000 loops=1)
14	Buckets: 16384 Batches: 1 Memory Usage: 656kB
15	→ Seq Scan on sailors s (cost=0.00..165.00 rows=9000 width=25) (actual time=0.008..0.982 rows=9000 loops=1)
16	→ Hash (cost=649.88..649.88 rows=7500 width=8) (actual time=6.678..6.680 rows=7569 loops=1)
17	Buckets: 8192 Batches: 1 Memory Usage: 360kB
18	→ Hash Join (cost=358.75..649.88 rows=7500 width=8) (actual time=2.345..5.875 rows=7569 loops=1)
19	Hash Cond: (r2.sid = s2.sid)
20	→ Hash Join (cost=81.25..352.68 rows=7500 width=4) (actual time=0.508..2.928 rows=7569 loops=1)
21	Hash Cond: (r2.bid = b2.bid)
22	→ Seq Scan on reserves r2 (cost=0.00..232.00 rows=15000 width=8) (actual time=0.012..0.687 rows=15000 loops=1)
23	→ Hash (cost=62.50..62.50 rows=1500 width=4) (actual time=0.487..0.487 rows=1500 loops=1)
24	Buckets: 2048 Batches: 1 Memory Usage: 69kB
25	→ Seq Scan on boat b2 (cost=0.00..62.50 rows=1500 width=4) (actual time=0.022..0.324 rows=1500 loops=1)
26	Filter: (color = 'red'::bpchar)
27	Rows Removed by Filter: 1500
28	→ Hash (cost=165.00..165.00 rows=9000 width=4) (actual time=1.777..1.778 rows=9000 loops=1)
29	Buckets: 16384 Batches: 1 Memory Usage: 445kB
30	→ Seq Scan on sailors s2 (cost=0.00..165.00 rows=9000 width=4) (actual time=0.012..0.748 rows=9000 loops=1)
31	Planning Time: 1.909 ms
32	Execution Time: 19.230 ms

Scenario 2 (B+ tree index):

1	explain analyze
2	select s.sname
3	from sailors s, reserves r, boat b

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Hash Semi Join (cost=1679.08..2276.43 rows=6250 width=21) (actual time=9.334..15.433 rows=7569 loops=1)		
2	Hash Cond: (s.sid = s2.sid)		
3	→ Hash Join (cost=577.60..1085.73 rows=7500 width=29) (actual time=3.064..7.891 rows=7569 loops=1)		
4	Hash Cond: (r.sid = s.sid)		
5	→ Hash Join (cost=143.81..632.25 rows=7500 width=4) (actual time=0.691..3.822 rows=7569 loops=1)		
6	Hash Cond: (r.bid = b.bid)		
7	→ Index Only Scan using reserves_pkey on reserves r (cost=0.29..449.29 rows=15000 width=8) (actual time=0.0...		
8	Heap Fetches: 0		
9	→ Hash (cost=124.78..124.78 rows=1500 width=4) (actual time=0.655..0.661 rows=1500 loops=1)		
10	Buckets: 2048 Batches: 1 Memory Usage: 69kB		
11	→ Index Scan using idx1 on boat b (cost=0.28..124.78 rows=1500 width=4) (actual time=0.013..0.490 rows=1...		
12	Filter: (color = 'red'::bpchar)		
13	Rows Removed by Filter: 1500		
14	→ Hash (cost=321.29..321.29 rows=9000 width=25) (actual time=2.322..2.324 rows=9000 loops=1)		
15	Buckets: 16384 Batches: 1 Memory Usage: 656kB		
16	→ Index Scan using idx4 on sailors s (cost=0.29..321.29 rows=9000 width=25) (actual time=0.008..1.269 rows=9...		
17	→ Hash (cost=1007.73..1007.73 rows=7500 width=8) (actual time=6.225..6.236 rows=7569 loops=1)		
18	Buckets: 8192 Batches: 1 Memory Usage: 360kB		
19	→ Hash Join (cost=499.60..1007.73 rows=7500 width=8) (actual time=2.122..5.497 rows=7569 loops=1)		
20	Hash Cond: (r2.sid = s2.sid)		
21	→ Hash Join (cost=143.81..632.25 rows=7500 width=4) (actual time=0.722..3.142 rows=7569 loops=1)		

22	Hash Cond: (r2.bid = b2.bid)		
23	→ Index Only Scan using reserves_pkey on reserves r2 (cost=0.29..449.29 rows=15000 width=8) (actual time...		
24	Heap Fetches: 0		
25	→ Hash (cost=124.78..124.78 rows=1500 width=4) (actual time=0.703..0.707 rows=1500 loops=1)		
26	Buckets: 2048 Batches: 1 Memory Usage: 69kB		
27	→ Index Scan using idx1 on boat b2 (cost=0.28..124.78 rows=1500 width=4) (actual time=0.010..0.553 ro...		
28	Filter: (color = 'red'::bpchar)		
29	Rows Removed by Filter: 1500		
30	→ Hash (cost=243.28..243.28 rows=9000 width=4) (actual time=1.351..1.353 rows=9000 loops=1)		
31	Buckets: 16384 Batches: 1 Memory Usage: 445kB		
32	→ Index Only Scan using idx4 on sailors s2 (cost=0.29..243.28 rows=9000 width=4) (actual time=0.012..0.54...		
33	Heap Fetches: 0		
34	Planning Time: 1.536 ms		
35	Execution Time: 15.891 ms		

Scenario 3 (hash index):

QUERY PLAN	
1	Hash Semi Join (cost=1102.38..1482.73 rows=6250 width=21) (actual time=9.920..16.285 rows=7569 loops=1)
2	Hash Cond: (s.sid = s2.sid)
3	-> Hash Join (cost=358.75..649.88 rows=7500 width=29) (actual time=2.425..7.494 rows=7569 loops=1)
4	Hash Cond: (r.sid = s.sid)
5	-> Hash Join (cost=81.25..352.68 rows=7500 width=4) (actual time=0.502..3.612 rows=7569 loops=1)
6	Hash Cond: (r.bid = b.bid)
7	-> Seq Scan on reserves r (cost=0.00..232.00 rows=15000 width=8) (actual time=0.017..0.931 rows=15000 loops=1)
8	-> Hash (cost=62.50..62.50 rows=1500 width=4) (actual time=0.471..0.476 rows=1500 loops=1)
9	Buckets: 2048 Batches: 1 Memory Usage: 69kB
10	-> Seq Scan on boat b (cost=0.00..62.50 rows=1500 width=4) (actual time=0.008..0.331 rows=1500 loops=1)
11	Filter: (color = 'red')::bpchar
12	Rows Removed by Filter: 1500
13	-> Hash (cost=165.00..165.00 rows=9000 width=25) (actual time=1.854..1.856 rows=9000 loops=1)
14	Buckets: 16384 Batches: 1 Memory Usage: 656kB
15	-> Seq Scan on sailors s (cost=0.00..165.00 rows=9000 width=25) (actual time=0.008..0.779 rows=9000 loops=1)
16	-> Hash (cost=649.88..649.88 rows=7500 width=8) (actual time=7.462..7.477 rows=7569 loops=1)
17	Buckets: 8192 Batches: 1 Memory Usage: 360kB
18	-> Hash Join (cost=358.75..649.88 rows=7500 width=8) (actual time=2.166..6.601 rows=7569 loops=1)
19	Hash Cond: (r2.sid = s2.sid)
20	-> Hash Join (cost=81.25..352.68 rows=7500 width=4) (actual time=0.410..3.363 rows=7569 loops=1)
21	Hash Cond: (r2.bid = b2.bid)
22	-> Seq Scan on reserves r2 (cost=0.00..232.00 rows=15000 width=8) (actual time=0.013..0.827 rows=15000 loops=1)
23	-> Hash (cost=62.50..62.50 rows=1500 width=4) (actual time=0.388..0.392 rows=1500 loops=1)
24	Buckets: 2048 Batches: 1 Memory Usage: 69kB
25	-> Seq Scan on boat b2 (cost=0.00..62.50 rows=1500 width=4) (actual time=0.007..0.262 rows=1500 loops=1)
26	Filter: (color = 'red')::bpchar
27	Rows Removed by Filter: 1500
28	-> Hash (cost=165.00..165.00 rows=9000 width=4) (actual time=1.708..1.711 rows=9000 loops=1)
29	Buckets: 16384 Batches: 1 Memory Usage: 445kB
30	-> Seq Scan on sailors s2 (cost=0.00..165.00 rows=9000 width=4) (actual time=0.008..0.657 rows=9000 loops=1)
31	Planning Time: 1.545 ms
32	Execution Time: 17.632 ms

Scenario 4 (mixed indexes):

1	Hash Semi Join (cost=1584.27..2129.62 rows=6250 width=21) (actual time=9.628..14.402 rows=7569 loops=1)
2	Hash Cond: (s.sid = s2.sid)
3	-> Hash Join (cost=556.20..1012.32 rows=7500 width=29) (actual time=3.390..7.233 rows=7569 loops=1)
4	Hash Cond: (r.sid = s.sid)
5	-> Hash Join (cost=122.41..558.84 rows=7500 width=4) (actual time=0.445..3.073 rows=7569 loops=1)
6	Hash Cond: (r.bid = b.bid)
7	-> Index Only Scan using idx17 on reserves r (cost=0.29..397.29 rows=15000 width=8) (actual time=0.013..1.133 rows=15000 loops=1)
8	Heap Fetches: 0
9	-> Hash (cost=103.38..103.38 rows=1500 width=4) (actual time=0.423..0.426 rows=1500 loops=1)
10	Buckets: 2048 Batches: 1 Memory Usage: 69kB
11	-> Bitmap Heap Scan on boat b (cost=59.63..103.38 rows=1500 width=4) (actual time=0.051..0.276 rows=1500 loops=1)
12	Recheck Cond: (color = 'red':bpchar)
13	Heap Blocks: exact=13
14	-> Bitmap Index Scan on idx11 (cost=0.00..59.25 rows=1500 width=0) (actual time=0.042..0.043 rows=1500 loops=1)
15	Index Cond: (color = 'red':bpchar)
16	-> Hash (cost=321.29..321.29 rows=9000 width=25) (actual time=2.893..2.896 rows=9000 loops=1)
17	Buckets: 16384 Batches: 1 Memory Usage: 656kB
18	-> Index Scan using idx4 on sailors s (cost=0.29..321.29 rows=9000 width=25) (actual time=0.009..1.515 rows=9000 loops=1)
19	-> Hash (cost=934.32..934.32 rows=7500 width=8) (actual time=6.191..6.202 rows=7569 loops=1)
20	Buckets: 8192 Batches: 1 Memory Usage: 360kB
21	-> Hash Join (cost=478.19..934.32 rows=7500 width=8) (actual time=2.101..5.457 rows=7569 loops=1)

22	Hash Cond: (r2.sid = s2.sid)
23	-> Hash Join (cost=122.41..558.84 rows=7500 width=4) (actual time=0.389..2.751 rows=7569 loops=1)
24	Hash Cond: (r2.bid = b2.bid)
25	-> Index Only Scan using idx17 on reserves r2 (cost=0.29..397.29 rows=15000 width=8) (actual time=0.013..1.002 rows=15000 loops=1)
26	Heap Fetches: 0
27	-> Hash (cost=103.38..103.38 rows=1500 width=4) (actual time=0.367..0.370 rows=1500 loops=1)
28	Buckets: 2048 Batches: 1 Memory Usage: 69kB
29	-> Bitmap Heap Scan on boat b2 (cost=59.63..103.38 rows=1500 width=4) (actual time=0.050..0.235 rows=1500 loops=1)
30	Recheck Cond: (color = 'red':bpchar)
31	Heap Blocks: exact=13
32	-> Bitmap Index Scan on idx11 (cost=0.00..59.25 rows=1500 width=0) (actual time=0.041..0.041 rows=1500 loops=1)
33	Index Cond: (color = 'red':bpchar)
34	-> Hash (cost=243.28..243.28 rows=9000 width=4) (actual time=1.663..1.666 rows=9000 loops=1)
35	Buckets: 16384 Batches: 1 Memory Usage: 445kB
36	-> Index Only Scan using idx4 on sailors s2 (cost=0.29..243.28 rows=9000 width=4) (actual time=0.014..0.680 rows=9000 loops=1)
37	Heap Fetches: 0
38	Planning Time: 1.529 ms
39	Execution Time: 15.039 ms

Schema 4

Query 10 without index

Query plan:

```
"Nested Loop (cost=228.29..231.72 rows=10 width=48) (actual time=0.162..0.163 rows=0 loops=1)"
"  -> HashAggregate (cost=228.00..228.10 rows=10 width=4) (actual time=0.162..0.163 rows=0 loops=1)"
"    Group Key: movie_cast.act_id"
"    Batches: 1 Memory Usage: 24kB"
"    -> Nested Loop (cost=0.29..227.98 rows=10 width=4) (actual time=0.162..0.162 rows=0 loops=1)"
"      -> Seq Scan on movie (cost=0.00..32.50 rows=1 width=4) (actual time=0.161..0.161 rows=0 loops=1)"
"        Filter: (mov_title = 'movie1'::bpchar)"
"        Rows Removed by Filter: 1000"
"      -> Index Only Scan using movie_cast_pkey on movie_cast (cost=0.29..195.38 rows=10 width=8) (never executed)"
"        Index Cond: (mov_id = movie.mov_id)"
"        Heap Fetches: 0"
"    -> Index Scan using actor_pkey on actor (cost=0.29..0.36 rows=1 width=48) (never executed)"
"      Index Cond: (act_id = movie_cast.act_id)"
"Planning Time: 0.322 ms"
"Execution Time: 0.202 ms"
```

RESULT CONCLUSION: Total time of 0.524 ms with no index.

Query 10 with B+ tree

Query plan:

```
"Nested Loop (cost=228.29..231.72 rows=10 width=48) (actual time=0.113..0.114 rows=0 loops=1)"
"  -> HashAggregate (cost=228.00..228.10 rows=10 width=4) (actual time=0.113..0.113 rows=0 loops=1)"
"    Group Key: movie_cast.act_id"
"    Batches: 1 Memory Usage: 24kB"
"    -> Nested Loop (cost=0.29..227.98 rows=10 width=4) (actual time=0.112..0.113 rows=0 loops=1)"
"      -> Seq Scan on movie (cost=0.00..32.50 rows=1 width=4) (actual time=0.112..0.112 rows=0 loops=1)"
"        Filter: (mov_title = 'movie1'::bpchar)"
"        Rows Removed by Filter: 1000"
"      -> Index Only Scan using movie_cast_pkey on movie_cast (cost=0.29..195.38 rows=10 width=8) (never executed)"
"        Index Cond: (mov_id = movie.mov_id)"
"        Heap Fetches: 0"
"    -> Index Scan using in1 on actor (cost=0.29..0.36 rows=1 width=48) (never executed)"
"      Index Cond: (act_id = movie_cast.act_id)"
"Planning Time: 0.313 ms"
"Execution Time: 0.134 ms"
```

RESULT CONCLUSION: Total time of 0.447 ms very slightly made faster when using a B+ tree index.

Query 10 with hash index

Query plan:

```
"Nested Loop (cost=66.48..67.47 rows=10 width=48) (actual time=0.129..0.130 rows=0 loops=1)"
"  -> HashAggregate (cost=66.48..66.58 rows=10 width=4) (actual time=0.129..0.129 rows=0 loops=1)"
"    Group Key: movie_cast.act_id"
"    Batches: 1 Memory Usage: 24kB"
"  -> Nested Loop (cost=4.08..66.45 rows=10 width=4) (actual time=0.128..0.128 rows=0 loops=1)"
"    -> Seq Scan on movie (cost=0.00..32.50 rows=1 width=4) (actual time=0.127..0.127 rows=0 loops=1)"
"      Filter: (mov_title = 'movie1'::bpchar)"
"      Rows Removed by Filter: 1000"
"    -> Bitmap Heap Scan on movie_cast (cost=4.08..33.85 rows=10 width=8) (never executed)"
"      Recheck Cond: (mov_id = movie.mov_id)"
"    -> Bitmap Index Scan on in2 (cost=0.00..4.08 rows=10 width=0) (never executed)"
"      Index Cond: (mov_id = movie.mov_id)"
"  -> Index Scan using in1 on actor (cost=0.00..0.09 rows=1 width=48) (never executed)"
"    Index Cond: (act_id = movie_cast.act_id)"
"Planning Time: 0.271 ms"
"Execution Time: 0.168 ms"
```


RESULT CONCLUSION: Total time 0.439 ms on par with the b+ tree speed.

Query 10 with mixed index

Query plan:

```
"Nested Loop (cost=66.76..70.19 rows=10 width=48) (actual time=0.128..0.128 rows=0
loops=1)"
"  -> HashAggregate (cost=66.48..66.58 rows=10 width=4) (actual time=0.127..0.128 rows=0
loops=1)"
"    Group Key: movie_cast.act_id"
"    Batches: 1 Memory Usage: 24kB"
"    -> Nested Loop (cost=4.08..66.45 rows=10 width=4) (actual time=0.126..0.126 rows=0
loops=1)"
"      -> Seq Scan on movie (cost=0.00..32.50 rows=1 width=4) (actual time=0.126..0.126
rows=0 loops=1)"
"        Filter: (mov_title = 'movie1'::bpchar)"
"        Rows Removed by Filter: 1000"
"      -> Bitmap Heap Scan on movie_cast (cost=4.08..33.85 rows=10 width=8) (never
executed)"
"        Recheck Cond: (mov_id = movie.mov_id)"
"      -> Bitmap Index Scan on in2 (cost=0.00..4.08 rows=10 width=0) (never
executed)"
"        Index Cond: (mov_id = movie.mov_id)"
"    -> Index Scan using in1 on actor (cost=0.29..0.36 rows=1 width=48) (never executed)"
"      Index Cond: (act_id = movie_cast.act_id)"
"Planning Time: 0.266 ms"
"Execution Time: 0.158 ms"
```

RESULT CONCLUSION: Total time of 0.424 ms giving the best result set by mixing both hash based and b+ index for this query.



Query 11 without index

Query plan:

```
"Nested Loop (cost=626.34..661.64 rows=100 width=42) (actual time=0.151..0.153 rows=0
loops=1)"
"  -> HashAggregate (cost=626.06..627.06 rows=100 width=4) (actual time=0.151..0.152
rows=0 loops=1)"
"    Group Key: movie_direction.dir_id"
"    Batches: 1 Memory Usage: 24kB"
"    -> Hash Semi Join (cost=453.46..625.81 rows=100 width=4) (actual time=0.150..0.151
rows=0 loops=1)"
"      Hash Cond: (movie_direction.mov_id = movie_cast.mov_id)"
"      -> Seq Scan on movie_direction (cost=0.00..144.99 rows=9999 width=8) (actual
time=0.010..0.010 rows=1 loops=1)"
"      -> Hash (cost=453.33..453.33 rows=10 width=4) (actual time=0.138..0.139 rows=0
loops=1)"
"        Buckets: 1024 Batches: 1 Memory Usage: 8kB"
"        -> Hash Semi Join (cost=242.99..453.33 rows=10 width=4) (actual
time=0.137..0.138 rows=0 loops=1)"
"          Hash Cond: (movie_cast.role = movie_cast_1.role)"
"          -> Seq Scan on movie_cast (cost=0.00..183.99 rows=9999 width=35) (actual
time=0.005..0.006 rows=1 loops=1)"
```

```

"          -> Hash (cost=242.86..242.86 rows=10 width=31) (actual time=0.130..0.131
rows=0 loops=1)"
"          Buckets: 1024 Batches: 1 Memory Usage: 8kB"
"          -> Hash Join (cost=32.51..242.86 rows=10 width=31) (actual
time=0.130..0.131 rows=0 loops=1)"
"          Hash Cond: (movie_cast_1.mov_id = movie.mov_id)"
"          -> Seq Scan on movie_cast movie_cast_1 (cost=0.00..183.99
rows=9999 width=35) (actual time=0.005..0.005 rows=1 loops=1)"
"          -> Hash (cost=32.50..32.50 rows=1 width=4) (actual
time=0.125..0.125 rows=0 loops=1)"
"          Buckets: 1024 Batches: 1 Memory Usage: 8kB"
"          -> Seq Scan on movie (cost=0.00..32.50 rows=1 width=4) (actual
time=0.124..0.124 rows=0 loops=1)"
"          Filter: (mov_title = 'movie'"
"2'::bpchar)"
"          Rows Removed by Filter: 1000"
" -> Index Scan using director_pkey on director (cost=0.28..0.35 rows=1 width=46) (never
executed)"
"      Index Cond: (dir_id = movie_direction.dir_id)"
"Planning Time: 0.387 ms"
"Execution Time: 0.181 ms"

```

RESULT CONCLUSION: Total time 0.568 ms without an index.

Query 11 with B+ tree

Query plan:

```

"Nested Loop (cost=450.21..485.52 rows=100 width=42) (actual time=0.129..0.131 rows=0
loops=1)"
" -> HashAggregate (cost=449.94..450.94 rows=100 width=4) (actual time=0.129..0.130
rows=0 loops=1)"
"      Group Key: movie_direction.dir_id"
"      Batches: 1 Memory Usage: 24kB"
"      -> Hash Semi Join (cost=277.34..449.69 rows=100 width=4) (actual time=0.128..0.129
rows=0 loops=1)"
"          Hash Cond: (movie_direction.mov_id = movie_cast.mov_id)"
"          -> Seq Scan on movie_direction (cost=0.00..144.99 rows=9999 width=8) (actual
time=0.007..0.008 rows=1 loops=1)"
"          -> Hash (cost=277.21..277.21 rows=10 width=4) (actual time=0.119..0.119 rows=0
loops=1)"
"          Buckets: 1024 Batches: 1 Memory Usage: 8kB"

```

```

"      -> Hash Semi Join (cost=66.86..277.21 rows=10 width=4) (actual
time=0.118..0.119 rows=0 loops=1)"
"      Hash Cond: (movie_cast.role = movie_cast_1.role)"
"      -> Seq Scan on movie_cast (cost=0.00..183.99 rows=9999 width=35) (actual
time=0.005..0.005 rows=1 loops=1)"
"      -> Hash (cost=66.74..66.74 rows=10 width=31) (actual time=0.112..0.112
rows=0 loops=1)"
"      Buckets: 1024 Batches: 1 Memory Usage: 8kB"
"      -> Nested Loop (cost=4.36..66.74 rows=10 width=31) (actual
time=0.111..0.112 rows=0 loops=1)"
"      -> Seq Scan on movie (cost=0.00..32.50 rows=1 width=4) (actual
time=0.111..0.111 rows=0 loops=1)"
"      Filter: (mov_title = 'movie'"
"2'::bpchar)"
"      Rows Removed by Filter: 1000"
"      -> Bitmap Heap Scan on movie_cast movie_cast_1 (cost=4.36..34.14
rows=10 width=35) (never executed)"
"      Recheck Cond: (mov_id = movie.mov_id)"
"      -> Bitmap Index Scan on in3 (cost=0.00..4.36 rows=10 width=0)
(never executed)"
"      Index Cond: (mov_id = movie.mov_id)"
" -> Index Scan using in1 on director (cost=0.28..0.35 rows=1 width=46) (never executed)"
"      Index Cond: (dir_id = movie_direction.dir_id)"
"Planning Time: 0.445 ms"
"Execution Time: 0.168 ms"

```

RESULT CONCLUSION: total time of 0.613 ms which is a significant increase from using no index.

Query 11 with hash index

Query plan:

```

"Nested Loop (cost=449.65..457.86 rows=100 width=42) (actual time=0.174..0.176 rows=0
loops=1)"
" -> HashAggregate (cost=449.65..450.65 rows=100 width=4) (actual time=0.174..0.175
rows=0 loops=1)"
"      Group Key: movie_direction.dir_id"
"      Batches: 1 Memory Usage: 24kB"
"      -> Hash Semi Join (cost=277.05..449.40 rows=100 width=4) (actual time=0.172..0.174
rows=0 loops=1)"
"      Hash Cond: (movie_direction.mov_id = movie_cast.mov_id)"

```

```

"      -> Seq Scan on movie_direction (cost=0.00..144.99 rows=9999 width=8) (actual
time=0.007..0.007 rows=1 loops=1)"
"      -> Hash (cost=276.93..276.93 rows=10 width=4) (actual time=0.163..0.164 rows=0
loops=1)"
"          Buckets: 1024 Batches: 1 Memory Usage: 8kB"
"      -> Hash Semi Join (cost=66.58..276.93 rows=10 width=4) (actual
time=0.163..0.164 rows=0 loops=1)"
"          Hash Cond: (movie_cast.role = movie_cast_1.role)"
"      -> Seq Scan on movie_cast (cost=0.00..183.99 rows=9999 width=35) (actual
time=0.005..0.005 rows=1 loops=1)"
"      -> Hash (cost=66.45..66.45 rows=10 width=31) (actual time=0.157..0.157
rows=0 loops=1)"
"          Buckets: 1024 Batches: 1 Memory Usage: 8kB"
"      -> Nested Loop (cost=4.08..66.45 rows=10 width=31) (actual
time=0.156..0.157 rows=0 loops=1)"
"          -> Seq Scan on movie (cost=0.00..32.50 rows=1 width=4) (actual
time=0.156..0.156 rows=0 loops=1)"
"              Filter: (mov_title = 'movie"
"2':::bpchar)"
"              Rows Removed by Filter: 1000"
"          -> Bitmap Heap Scan on movie_cast movie_cast_1 (cost=4.08..33.85
rows=10 width=35) (never executed)"
"              Recheck Cond: (mov_id = movie.mov_id)"
"          -> Bitmap Index Scan on in3 (cost=0.00..4.08 rows=10 width=0)
(never executed)"
"              Index Cond: (mov_id = movie.mov_id)"
"      -> Index Scan using in1 on director (cost=0.00..0.08 rows=1 width=46) (never executed)"
"          Index Cond: (dir_id = movie_direction.dir_id)"
"Planning Time: 0.476 ms"
"Execution Time: 0.209 ms"

```

RESULT CONCLUSION: Total time of 0.685 ms an even greater decrease in performance from using b+ index.

Query 11 with mixed index

Query plan:

```

"Nested Loop (cost=449.93..485.23 rows=100 width=42) (actual time=0.270..0.273 rows=0
loops=1)"
"  -> HashAggregate (cost=449.65..450.65 rows=100 width=4) (actual time=0.269..0.272
rows=0 loops=1)"
"      Group Key: movie_direction.dir_id"

```



```

"      Batches: 1 Memory Usage: 24kB"
"      -> Hash Semi Join (cost=277.05..449.40 rows=100 width=4) (actual time=0.267..0.270
rows=0 loops=1)"
"          Hash Cond: (movie_direction.mov_id = movie_cast.mov_id)"
"          -> Seq Scan on movie_direction (cost=0.00..144.99 rows=9999 width=8) (actual
time=0.012..0.012 rows=1 loops=1)"
"          -> Hash (cost=276.93..276.93 rows=10 width=4) (actual time=0.251..0.254 rows=0
loops=1)"
"              Buckets: 1024 Batches: 1 Memory Usage: 8kB"
"              -> Hash Semi Join (cost=66.58..276.93 rows=10 width=4) (actual
time=0.251..0.253 rows=0 loops=1)"
"                  Hash Cond: (movie_cast.role = movie_cast_1.role)"
"                  -> Seq Scan on movie_cast (cost=0.00..183.99 rows=9999 width=35) (actual
time=0.008..0.008 rows=1 loops=1)"
"                  -> Hash (cost=66.45..66.45 rows=10 width=31) (actual time=0.241..0.242
rows=0 loops=1)"
"                      Buckets: 1024 Batches: 1 Memory Usage: 8kB"
"                      -> Nested Loop (cost=4.08..66.45 rows=10 width=31) (actual
time=0.240..0.241 rows=0 loops=1)"
"                          -> Seq Scan on movie (cost=0.00..32.50 rows=1 width=4) (actual
time=0.240..0.240 rows=0 loops=1)"
"                              Filter: (mov_title = 'movie"
"2'::bpchar)"
"                                  Rows Removed by Filter: 1000"
"                                  -> Bitmap Heap Scan on movie_cast movie_cast_1 (cost=4.08..33.85
rows=10 width=35) (never executed)"
"                                      Recheck Cond: (mov_id = movie.mov_id)"
"                                      -> Bitmap Index Scan on in3 (cost=0.00..4.08 rows=10 width=0)
(never executed)"
"                                          Index Cond: (mov_id = movie.mov_id)"
"      -> Index Scan using in1 on director (cost=0.28..0.35 rows=1 width=46) (never executed)"
"          Index Cond: (dir_id = movie_direction.dir_id)"
"Planning Time: 0.446 ms"
"Execution Time: 0.323 ms"

```

RESULT CONCLUSION: Total time 0.769 ms which shows that using mixed index provides a significant decrease in performance on this partial query.



Query 12 without index

Query plan :

```
"Index Scan using movie_pkey on movie (cost=53.65..61.67 rows=1 width=51) (actual
time=0.182..0.183 rows=0 loops=1)"
"  Index Cond: (mov_id = $1)"
"  InitPlan 2 (returns $1)"
"    -> Index Only Scan using movie_direction_pkey on movie_direction (cost=49.28..53.37
rows=5 width=4) (actual time=0.180..0.180 rows=0 loops=1)"
"      Index Cond: (dir_id = $0)"
"      Heap Fetches: 0"
"      InitPlan 1 (returns $0)"
"        -> Seq Scan on director (cost=0.00..49.00 rows=1 width=4) (actual time=0.177..0.177
rows=0 loops=1)"
"          Filter: ((dir_fname = 'actor1'::bpchar) AND (dir_lname = 'actor1'::bpchar))"
"          Rows Removed by Filter: 2000"
"Planning Time: 0.110 ms"
"Execution Time: 0.200 ms"
```

RESULT CONCLUSION: Total time of 0.310 ms without an index.

Query 12 with b+ tree

Query plan :

```
"Index Scan using in1 on movie (cost=53.65..61.67 rows=1 width=51) (actual
time=0.185..0.186 rows=0 loops=1)"
"  Index Cond: (mov_id = $1)"
"  InitPlan 2 (returns $1)"
"    -> Index Only Scan using movie_direction_pkey on movie_direction (cost=49.28..53.37
rows=5 width=4) (actual time=0.183..0.184 rows=0 loops=1)"
"      Index Cond: (dir_id = $0)"
"      Heap Fetches: 0"
"      InitPlan 1 (returns $0)"
"        -> Seq Scan on director (cost=0.00..49.00 rows=1 width=4) (actual time=0.180..0.180
rows=0 loops=1)"
"          Filter: ((dir_fname = 'actor1'::bpchar) AND (dir_lname = 'actor1'::bpchar))"
"          Rows Removed by Filter: 2000"
"Planning Time: 0.125 ms"
"Execution Time: 0.203 ms"
```

RESULT CONCLUSION: Total time of 0.328 ms not a significant change from using no index.

Query 12 with hash index

Query plan :

```
"Index Scan using in1 on movie (cost=53.37..61.39 rows=1 width=51) (actual
time=0.220..0.221 rows=0 loops=1)"
"  Index Cond: (mov_id = $1)"
"    InitPlan 2 (returns $1)"
"      -> Index Only Scan using movie_direction_pkey on movie_direction (cost=49.28..53.37
rows=5 width=4) (actual time=0.218..0.219 rows=0 loops=1)"
"        Index Cond: (dir_id = $0)"
"        Heap Fetches: 0"
"        InitPlan 1 (returns $0)"
"          -> Seq Scan on director (cost=0.00..49.00 rows=1 width=4) (actual time=0.199..0.199
rows=0 loops=1)"
"            Filter: ((dir_fname = 'actor1'::bpchar) AND (dir_lname = 'actor1'::bpchar))"
"            Rows Removed by Filter: 2000"
"Planning Time: 0.143 ms"
"Execution Time: 0.259 ms"
```

RESULT CONCLUSION: Total time of 0.402 ms a significant performance decrease using a hash index.

Query 12 with mixed index

Query plan :

```
"Index Scan using in1 on movie (cost=53.37..61.39 rows=1 width=51) (actual
time=0.243..0.244 rows=0 loops=1)"
"  Index Cond: (mov_id = $1)"
"  InitPlan 2 (returns $1)"
"    -> Index Only Scan using movie_direction_pkey on movie_direction (cost=49.28..53.37
rows=5 width=4) (actual time=0.241..0.241 rows=0 loops=1)"
"      Index Cond: (dir_id = $0)"
"      Heap Fetches: 0"
"      InitPlan 1 (returns $0)"
"        -> Seq Scan on director (cost=0.00..49.00 rows=1 width=4) (actual time=0.236..0.236
rows=0 loops=1)"
"          Filter: ((dir_fname = 'actor1'::bpchar) AND (dir_lname = 'actor1'::bpchar))"
"          Rows Removed by Filter: 2000"
"Planning Time: 0.116 ms"
"Execution Time: 0.283 ms"
```

RESULT CONCLUSION: Total time 0.399 ms using mixed index gave similar results to using a hash index both being lower performance wise from no index or b+ tree index.

